

AFNAN PK

Agentic AI Engineer

afnaan.codes@gmail.com | +971 050 732 7220 | linkedin.com/in/afnan-pk | github.com/afnaann

afnaann.vercel.app | Dubai, UAE

SUMMARY

Agentic AI Engineer with 2+ years of experience building and deploying AI-powered software, including production LLM, RAG, and multi-agent systems. Experienced with LangGraph, LangChain, OpenAI API, Qdrant, LLM evaluation, MCP, asynchronous backend systems, and production AI optimization. Built customer-facing AI systems handling hundreds of daily interactions, with measurable improvements in response latency, LLM cost, and production reliability.

SKILLS

AI & LLM Engineering — LangGraph, LangChain, OpenAI API, RAG, Hybrid Retrieval, Qdrant, Vector Embeddings, Multi-Agent Systems, MCP, Prompt Engineering, LLM Evaluation, LangSmith, Semantic Caching, Model Routing, Human-in-the-Loop Workflows

Backend & Distributed Systems — Python, Node.js, Django, Express.js, REST APIs, Redis, BullMQ, Celery, asyncio, Microservices, Webhooks, Asynchronous Processing

Data, Cloud & Infrastructure — Apache Airflow, PostgreSQL, MongoDB, AWS S3, AWS Redshift, Azure, Docker, CI/CD, GitHub

Frontend — React, Next.js, TypeScript, Redux Toolkit, Tailwind CSS

EXPERIENCE

Wiral AI

07/2025 – 08/2026 | Doha, Qatar

Agentic AI Engineer

- Engineered a production multi-channel AI customer engagement platform currently handling approximately **400–600 daily customer interactions** across sales and support workflows.
- Redesigned the existing **single-agent architecture** into a **LangGraph-based multi-agent system** with dedicated Inquiry, Booking, and Back Office agents, separating customer-facing RAG, transactional booking, and asynchronous lead-processing workflows.
- Reduced user-facing response latency by approximately **30%** by moving non-critical post-response workloads, including attribute extraction and lead classification, into asynchronous back-office processing.
- Engineered the **RAG pipeline** for knowledge-base ingestion and web content extraction, implementing **hybrid keyword and vector retrieval** with Qdrant and tenant-level data isolation.
- Built asynchronous backend processing with **Node.js, Redis, and BullMQ** for webhook ingestion, sequential message processing, retries, and automated follow-ups; implemented **Redis-based distributed locking** to prevent concurrent conversation-state corruption and duplicate responses.
- Integrated **MCP with Zoho CRM** to enable AI workflows to retrieve and manage lead information, with **human-in-the-loop escalation** for complex conversations requiring manual intervention.
- Implemented **LangSmith tracing and evaluation workflows** using test cases and manual trace review to identify response and workflow failures, iterating prompts, orchestration, and model selection; reduced measured LLM costs by **40–50%** through semantic caching and cost-aware model routing.
- Deployed and maintained production AI services on **Azure VMs** using **Docker** across development, UAT, and production environments.

WebMavericks Softcoders

01/2025 – 07/2025

AI Full Stack Engineer

- Built and optimized **50+ Apache Airflow ETL pipelines** using Python, integrating data processing and **LLM-assisted data cleaning** with AWS S3 and Redshift for analytics and reporting workloads.

- Optimized a Django-based QR scan-and-track system for a mining logistics platform, reducing checkpoint processing time from approximately **12 seconds to under 3 seconds**.
- Developed REST APIs within a Django/PostgreSQL operations platform, implementing role-based access control for secure multi-role workflows.

Bridgeon Solutions03/2024 – 01/2025

Full Stack Developer

- Co-architected a **six-service microservices-based e-commerce platform** covering user, product, delivery, and administrative workflows, using **RabbitMQ** for asynchronous inter-service communication.
- Integrated third-party APIs and internal services to support cross-service workflows and data exchange across the platform.
- Built **CI/CD workflows** with automated test execution and release automation to improve deployment consistency across the microservices architecture.

EDUCATION

Bachelor of Computer Applications2024 – 2027 | Delhi, India

IGNOU UNIVERSITY [🔗](#)

LANGUAGES

English — Expert

Hindi — Proficient

Arabic — Basic

Malayalam — Native